

DOMAIN: FinTech

CHOSEN COMPANY: KreditBee

DATE: 08-04-2026

PROJECT BY: Allan Cheerakunnil Alex

Business Understand & DBMS Foundation

Group Members :

1. Allan Cheerakunnil Alex
2. Babitha Tomy
3. Faisal Khan
4. Kartihik Reddy
5. Madan Dahiphale
6. Mansahib Singh

Domain : Fintech Industry

Company : Kredit Bee

Date : 11/04/2026

CASE STUDY

DA03: Financial Risk and Performance Analysis in Consumer Lending

The FinTech industry has emerged as a dynamic and transformative force in the financial sector, integrating technology with financial services to enhance efficiency, accessibility, and customer experience. In today's era, FinTech is crucial for democratizing financial services, providing underserved populations with access to banking, credit, and investment opportunities. It fosters innovation through digital payment systems, peer-to-peer lending platforms, and blockchain technology, revolutionizing traditional banking practices.

Data scientists play a pivotal role in the growth of FinTech by leveraging advanced analytics and machine learning to improve risk assessment, fraud detection, and personalized financial services. Their expertise enables FinTech companies to analyze vast amounts of financial data, uncovering insights that drive strategic decision-making, optimize operations, and enhance customer satisfaction. By harnessing the power of data, data scientists help FinTech firms stay competitive, innovate continuously, and contribute to a more inclusive and efficient financial ecosystem.

Product Dissection

1. Choose a leading platform from a domain related to the FinTech industry. Justify your selection by discussing the platform's popularity, impact, and relevance in its industry.

Answer:

Company Chosen: KreditBee

- **Popularity:** KreditBee is one of India's leading digital lending platforms, boasting over 100 million app downloads and a customer base exceeding 7 million active users. It has achieved massive scale with a credit disbursement tally crossing ₹30,000 crore.
- **Impact:** KreditBee disrupted the traditional banking sector by introducing a fully automated digital lending model. It bypassed legacy banking constraints by utilizing alternative data points (device data, location, SMS parsing) to underwrite loans for users who lacked traditional CIBIL scores, reducing loan disbursal times from days to a mere 10 minutes.
- **Relevance:** The Indian middle class and young professional demographic often face a "credit gap"—they need short-term liquidity for emergencies or purchases but lack the assets to secure traditional bank loans. KreditBee captures this "missing middle" market by providing unsecured personal loans ranging from ₹6,000 to ₹10

Lakhs, making formal credit accessible and driving financial inclusion in India.

2. Research and list the core features and functionalities of the selected platform. Describe how these features contribute to the platform's success and user engagement.

Answer:

- **Feature 1: Tiered Personal Loans (Flexi & Salaried):** KreditBee offers highly customized loan products—ranging from small-ticket "Flexi Loans" for immediate needs to larger "Salaried Loans". *Impact:* This allows them to acquire a user at a low ticket size, build their credit history, and later upsell larger loans, directly increasing Customer Lifetime Value (LTV).
- **Feature 2: 100% Digital KYC & Automated Underwriting:** The platform uses OCR (Optical Character Recognition) for instant document verification and API integrations with credit bureaus. *Impact:* It drastically lowers Customer Acquisition Cost (CAC) by removing human intervention and operational bottlenecks, ensuring instant approvals.
- **Feature 3: KreditBee UPI Integration:** Users can link their bank accounts to a KreditBee UPI ID to make secure online and offline payments. *Impact:* This transitions KreditBee from a purely transactional loan app into an everyday utility app, massively increasing Daily Active Users (DAU) and app engagement.

- **Feature 4: Seamless Repayment Ecosystem:** The app features integrated payment gateways and auto-debit setups for EMI collections. *Impact:* This reduces friction during the repayment cycle, directly contributing to their high collection efficiency (typically 94-97%).

3. Identify the real-world problems that the platform aims to solve. Discuss how the platform addresses these problems through its features and functionalities.

Answer:

- **Problem 1: Exclusion of "New to Credit" (NTC) Borrowers:**
Traditional banks reject young professionals because they lack a credit history.

Solution: KreditBee uses alternative data algorithms and AI-based risk profiling to underwrite young professionals and self-employed individuals without a rigid CIBIL footprint.
- **Problem 2: High Turn-Around-Time (TAT) in Emergencies:**
Applying for a bank loan requires physical branch visits, heavy paperwork, and waiting days for approval, which is useless during a medical emergency.

Solution: KreditBee's 100% online platform ensures eligibility checking, approval, and bank disbursement happen in under 10 minutes.
- **Problem 3: Exploitation by Informal Moneylenders:** Without access to formal banking, the working-class population relies on local loan sharks who charge 100-200% interest rates.

Solution: KreditBee provides a transparent, RBI-compliant alternative with clear processing fees and regulated APRs ranging from 17% to 50%.

Schema Design

4. Based on the features and functionalities you have identified, design a schema that reflects the platform's data structure. Define the key entities, attributes, and relationships that underpin these features.

Key Entities and Attributes

1. User_Profile Table

Attribute	Type	Description
userID	Primary Key (UUID)	Unique user identifier for the user
mobileNumber	String	Registered Phone Number

monthlyIncome	Decimal	Users declared or verified income
cibilScore	Integer	Credit Score (Null if New-to-Credit)
kycVerified	Boolean	Aadhaar/PAN verification status

2. Loan_Product Table

Attribute	Type	Description
productID	Primary Key	Unique user identifier
productName	String	e.g., Flexi Personal Loan
minAmount	Decimal	Minimum disbursal (e.g., ₹6000)
maxAmount	Decimal	Maximum disbursal allowed

baseInterestRate	Decimal	Base annual interest rate
------------------	---------	---------------------------

3. Loan_Application Table

Attribute	Type	Description
applicationID	Primary Key	Unique application identifier
userID	Foreign Key	Links to User_Profile
productID	Foreign Key	Minimum disbursal (e.g., ₹6000)
requestedAmount	Decimal	Amount requested by the user
riskTier	String	Risk category (A, B, C, High)
Status	String	Approved, Rejected, Disbursed, Defaulted

4. Repayment_Schedule Table

Attribute	Type	Description
EmiID	Primary Key	Unique EMI identifier

applicationID	Foreign Key	Link to Loan_Application
dueDate	Date	Payment due date
emiAmount	Decimal	Monthly installment amount
paymentStatus	String	Pending, Paid, Late, Default

5. Relationship Mapping Table


Relationship	Description
One-to-Many	One User_Profile → Multiple Loan_Applications
One-to-Many	One Loan_Product → Multiple Loan_Applications

One-to-Many	One Loan_Application → Multiple Repayment_Schedule entries
-------------	--

This schema design captures the essential data interactions and relationships within KreditBee, supporting its features and functionalities efficiently. It ensures scalability, performance, and a personalized user experience, contributing to the platform's overall success.

5. Utilise tools like the Miro platform or similar applications to create an illustrative Entity-Relationship (ER) diagram. This diagram should vividly depict the entities, attributes, and relationships present within your schema design.

Answer:



Revenue and Profit Growth Strategies

6. Question: After completing the product dissection and schema design steps for the chosen platform, conduct a comprehensive case study on the above chosen industry. Your goal is to identify and propose strategies to increase the profit of the industry by at least 25%

Analysing KreditBee’s Current Status

To understand KreditBee’s current position and build a strategy for 25% profit growth, we analyse:

- Current financial performance
 - Revenue sources
 - Expense structure
 - Customer acquisition and retention
-

Category-wise Analysis

1. Current Financial Data

Category	Details
Collect Financial Data	
Revenue	KreditBee earns from interest income, processing fees, penalties, and partner commissions. In FY25, service revenue was ₹486 crore and other income ₹103.8 crore (22% growth)

Expenses	Major costs include marketing, employee salaries, and operational expenses like collections and tech infrastructure.
Analyse Current Profit	
Profit Calculation	Profit = Revenue – Expenses
Profit Trends	KreditBee improved profit significantly by reducing costs like ad spend and employee expenses.

2. Sources of Revenue

Category	Details
Identify Revenue Streams	
Interest Income	Main revenue source from loans given to customers
Processing Fees	One-time fee charged during loan disbursement
Late Fees / Penalties	Charges for missed EMIs

Cross-Selling	Income from insurance, credit cards, and partner services
Other Services	Credit assessment and additional service fees
Analyse Revenue Contribution	
Segmentation	Revenue is divided into interest + fees + partner income
Key Insight	Strong diversified model reduces dependency on a single source

3. Sources of Expenses

Category	Details
Identify Major Expenses	
Marketing & Advertising	Major expense used to acquire new users
Technology Costs	Maintaining app, AI models, and infrastructure

Employee Salaries	Salaries for operations, tech, and support teams
Operational Costs	Collections, fraud checks, and backend processes
Analyse Expense Distribution	
Segmentation	Marketing is the highest cost contributor
Cost Efficiency	KreditBee spends ₹0.61 to earn ₹1 — improving this ratio increases profit

4. Customer Acquisition & Retention

Category	Details
Analyse Customer Acquisition Channels	
Channels	Digital ads, app installs, partnerships, referrals
Effectiveness	High marketing spend indicates dependency on paid acquisition

Understand Customer Behaviour & Retention	
Customer Data	Users are mostly Tier-2/3 customers with repeat borrowing needs
Retention Rates	Repeat users contribute higher revenue with lower cost
Churn Analysis	Drop-offs occur due to pricing, trust, or repayment issues

Conclusion of Analysis

By analysing revenue, expenses, and customer behaviour, we identify:

- Revenue can be increased by **expanding loans and improving pricing**
 - Costs can be reduced through **automation and better marketing strategies**
 - Retention can improve **lifetime value and profitability**
-

II. Focus Areas for Increasing KreditBee's Profit by 25%

To achieve 25% profit growth, KreditBee should focus on key business areas:

Strategy Table

Category	Focus Area	Measures
----------	------------	----------

Internal Management ~8%	Operational Efficiency ~6%	Automate loan processing, fraud detection, and collections using AI.Optimize internal processes to handle more loans with fewer resources.Control costs by reducing unnecessary ad spend and improving CAC.
	Employee Productivity ~2%	Train teams in data-driven decision making.Use performance metrics like approval rate, default rate, and efficiency tracking.
Financial Growth ~10–15%	Loan Book Expansion ~10–15%	Increase loan disbursement in Tier-2/3 markets.Launch new products like BNPL and credit cards.Expand partnerships to acquire more users.
	Pricing Optimization ~5%	Use ML-based risk scoring to adjust interest rates.Offer premium services like instant disbursal for extra fees.Improve profit margins through smarter pricing.
Cost Management ~5–10%	Expense Reduction ~5%	Reduce marketing cost using referral programs.Automate customer support and

		collections.Negotiate better funding and vendor terms.
	Unit Economics Improvement ~5%	Improve cost-to-revenue ratio (₹0.61 → ₹0.55).Use fraud detection to reduce bad loans.
Customer Growth ~5–10%	Customer Acquisition ~5%	Focus on underserved regions using local-language campaigns.Use partnerships instead of expensive ads.
	Customer Retention ~5%	Introduce loyalty programs and rewards.Encourage repeat loans through pre-approved offers.Use reminders to reduce missed payments.

III. Defining Strategies

Category	Details
Optimize Expenses	Reduce operational costs through automation, better vendor deals, and efficient marketing strategies.
Enhance Revenue Streams	Increase loan volume, introduce new financial products, and improve pricing strategies using data.
Improve Customer Retention	Provide personalized offers, loyalty rewards, and better user experience to increase repeat usage.

Final Conclusion

By following an inside-out approach:

- First optimize **expenses**
- Then improve **customer retention**
- Finally scale **revenue streams**

KreditBee can achieve **25% profit growth** in a structured and sustainable way.

Guesstimate Questions

1. What will be the percentage increase in global FinTech investments over the next five years?

Compound Annual Growth Rate (CAGR) model. Despite recent funding winters, the foundational shift to digital finance (especially in APAC and Africa) guarantees steady long-term growth.

Step	Assumption	Calculation
1	Current Global FinTech Investment (Base Year)	Assume \$100 Billion (for mathematical simplicity)
2	Expected Annual Growth Rate (CAGR)	Assume 15% (industry standard projection)

3	Growth Multiplier	$1 + 0.15 = \mathbf{1.15}$
4	Value after 5 Years	$\$100\text{B} \times (1.15)^5 \approx \mathbf{\$201 \text{ Billion}}$
5	Percentage Increase	$((\$201\text{B} - \$100\text{B}) / \$100\text{B}) * 100 = \mathbf{101\%}$

Final Answer: We estimate a roughly **100% to 105%** increase in global FinTech investments over the next five years, effectively doubling the current market size.

2. How many people will adopt digital banking services in developing countries over the next decade?

We have used India as a proxy for the developing world, utilizing a top-down population funnel based on smartphone penetration and internet access.

Step	Assumption	Calculation	Result
1	Proxy Population (India)	Current population	1.4 Billion

2	Adult Population (18+)	$1.4B \times 65\%$	~910 Million
3	Estimated Smartphone/Internet Users by 2030	$910M \times 80\%$ penetration	~728 Million
4	Currently Using Digital Banking (e.g., UPI/Apps)	$728M \times 50\%$	~364 Million
5	Remaining Unbanked/Cash-heavy Adults	$728M - 364M$	~364 Million
6	Expected Adoption Rate over 10 Years	$364M \times 75\%$ adoption	~273 Million

Final Answer: In India alone, we estimate **270 to 280 million** new users will adopt digital banking. Scaled across other developing nations (like Brazil, Indonesia, Nigeria), the global figure will likely exceed **600 to 700 million** new adopters over the next decade.

3. What percentage of small and medium-sized enterprises (SMEs) will use FinTech solutions for their financial needs by 2026?

SMEs are rapidly digitizing because traditional banks require heavy collateral that SMEs lack. FinTechs (like BharatPe or KreditBee Business) offer cash-flow-based lending.

Step	Assumption	Calculation	Result
1	Total MSMEs in India	Official Government Estimates	~63 Million
2	Digitally Addressable SMEs	$63M \times 60\%$ (internet access/GST registered)	~37.8 Million
3	Current FinTech Adoption Rate	$37.8M \times 40\%$	~15.1 Million
4	Projected Adoption Growth by 2025	Increase to 70% of addressable market $37.8M \times 70\%$	~26.46 Million
5	Total SMEs Using FinTech by 2026	Projected Target	~26.4 Million
6	Final Percentage (Total Market)	$(26.4M / 63M) \times 100$	~42%

Final Answer: We estimate that approximately **42% of all SMEs** and

up to **70% of digitally addressable SMEs** will heavily rely on FinTech solutions by 2026.

4. What will be the average transaction value of mobile payments in the next three years?

Mobile payments started with micro-transactions (buying tea/groceries) but are now moving to high-ticket items (rent, electronics, loan EMI) due to increased trust.

Step	Assumption	Calculation	Result
1	Current Avg. Mobile Payment Value (P2M/P2P)	Baseline Assumption	₹1,600
2	Inflation Rate	Annualized	~6%
3	Digital Trust Factor (Larger purchases)	Additional growth	~4%
4	Total Annual Growth Rate	6% + 4%	10%
5	Growth Multiplier	1 + 0.10	1.10

6	Value after 3 Years	$\text{₹}1,600 \times (1.10)^3$	₹2,129
---	---------------------	---------------------------------	---------------

Final Answer: Driven by inflation and higher consumer confidence in digital limits, the average transaction value is expected to grow to approximately **₹2,100 - ₹2,150** in the next three years.

5. How much will blockchain technology reduce the costs of cross-border transactions in the next five years?

Traditional remittances (like Western Union or SWIFT) involve multiple intermediary banks, Nostro/Vostro accounts, and forex margins. Blockchain enables direct peer-to-peer settlement.

Step	Assumption	Calculation	Result
1	Current Avg. Cross-Border Transaction Fee	World Bank Data	~6.0%
2	Target Blockchain Settlement Fee	Direct network fees (e.g., Ripple/Stellar)	~1.5%
3	Total Fee Reduction	6.0% - 1.5%	4.5% absolute reduction

4	Percentage Reduction in Cost	$(4.5\% / 6.0\%) * 100$	75% relative reduction
----------	------------------------------	-------------------------	-------------------------------

Final Answer: Blockchain technology has the potential to reduce cross-border transaction costs by roughly **75%**, bringing the average fee down from 6% to around 1.5%.

Cohort Analysis

Question 1 : How would you perform a cohort analysis to calculate the monthly retention rate for users who signed up in different months?

1. Define Cohort

Users are grouped into cohorts based on their loan issue date (`issue_d`). Since the dataset does not contain sufficient variation at the monthly level, cohorts are defined at a year level (`cohort_year`) as a proxy for signup time.

2. Define User Activity (Proxy)

As transaction-level activity is not available, user engagement is approximated using repayment behaviour:

- A user is considered active if:
 - `total_pymnt > 0`, and
 - `loan_condition = Good Loan`

This indicates that the user is still financially engaged and has not defaulted.

3. Calculate Retention

For each cohort:

- Total Users = Number of unique users in the cohort
- Active Users = Users satisfying the activity condition
- Retention Rate = Active Users / Total Users

Cohort Year	Total Users	Active Users	Retention Rate
2007	139	95	0.683453
2008	516	417	0.808140
2009	1173	1003	0.855072
2010	2822	2400	0.850461
2011	4747	4031	0.849168
2012	11777	9899	0.840537
2013	29698	25575	0.861169
2014	19128	17616	0.920954

4. Compare Cohorts (Insights)

- Retention improves from 2007 (68.3%) → 2009 (85.5%), indicating better early cohort performance.
- Retention remains relatively stable between 2010–2012 (~84–85%).
- A noticeable increase is observed in later cohorts, with 2014 reaching 92.1%, the highest retention.

Overall, newer cohorts show improved retention, suggesting improvements in user quality, product experience, or risk management over time.

Q2. If you find that retention drops significantly after the first month, what are some possible reasons for this behaviour, and how can the company address this drop-off?

1. Possible Reasons for Drop in Retention

A significant drop in retention after the first month, our case year can be attributed to the following factors:

a. User Quality Issues

- Users with high debt-to-income (DTI) or lower credit grades may struggle to continue repayments.
- This leads to early disengagement (default or inactivity).

b. Pricing & Financial Burden

- Higher interest rates increase repayment pressure.
 - Users may drop off after initial interaction due to affordability issues.
-

c. Lack of Engagement / Awareness

- Users may not clearly understand repayment schedules.
 - Absence of reminders or guidance leads to missed payments.
-

d. Poor Targeting

- Loans may be issued to financially unstable users (low income, low employment stability).
 - These users are less likely to sustain long-term engagement.
-

2. How to Address the Drop-Off

a. Improve Risk-Based Targeting

- Prioritize users with:
 - Lower DTI
 - Higher credit grades
 - Strengthen credit evaluation models.
-

b. Optimize Pricing Strategy

- Offer flexible repayment options
 - Adjust interest rates based on risk segments
-

c. Enhance User Engagement

- Send EMI reminders and notifications
 - Provide clear onboarding about repayment schedules
-

d. Introduce Incentives

- Cashback or rewards for on-time or early payments
 - Discounts for completing initial repayments
-

3. Data-Driven Validation (Advanced Insight)

The hypotheses can be validated by segmenting users based on:

- `dti` (debt-to-income ratio)
- `grade` (credit quality)
- `income_category`

and comparing retention rates across these segments.

A/B Test (2012 vs 2013 Cohorts)

Question 1:

How would you structure an A/B test to measure the impact of these notification designs on loan acceptance rates?

Step 1 — Define the Business Problem

Question:

Does offering a **reduced interest rate** encourage more users to accept loans without increasing default risk?

Groups:

Group	Year	Strategy
A (Control)	2012	Default interest rate
B (Treatment)	2013	Reduced interest rate

Step 2 — State the Hypotheses

- **H₀ (Null):** Reduced interest rates **do not change** loan acceptance or repayment behavior.
- **H₁ (Alternative):** Reduced interest rates **increase loan acceptance** without any risk.

Step 3 — Select Evaluation Metrics

Use fields available in your dataset.

Metric	Column(s)	Why It Matters
Loan Acceptance Rate	loan_condition	Core conversion outcome
Average Loan Amount	loan_amount	Borrowing confidence
Repayment Behaviour	total_rec_prncp, total_pymnt	Repayment quality
Default Risk	recoveries	Portfolio risk indicator

Step 4 — Prepare the Groups (Filter the Data)

Create two subsets:

- **Group A:** records where year = 2012
- **Group B:** records where year = 2013

This simulates **control vs treatment cohorts**.

Step 5 — Compute Metrics for Each Group

For **each group**, calculate:

1. **Loan Acceptance Rate**
= Good Loans / Total Loans
(derived from loan_condition)
 2. **Average Loan Amount**
= mean(loan_amount)
 3. **Repayment Behaviour**
= average repayment level using total_rec_prncp or total_pymnt
 4. **Default Risk**
= average or total recoveries
-

Step 6 — Build the Final Comparison Table

This is your **core A/B test output**.

Group	Loan Acceptance Rate	Avg Loan Amount	Repayment Behaviour	Default Risk
Group A (2012)	0.841	13535.92	11466.81	138.90
Group B (2013)	0.861	14740.59	10666.37	119.25

This table visually shows **which strategy performs better**.

Step 7 — Interpret the Results

Evaluate directional change between Group A (2012) and Group B (2013).

Observation	Interpretation
Loan Acceptance Rate increased (0.8405 → 0.8612)	Borrowers are slightly more willing to take loans in Group B. This suggests the lending conditions in 2013 improved loan uptake.
Average Loan Amount increased (13,535 → 14,740)	Borrowers are requesting larger loans, indicating greater borrowing confidence or higher credit availability.
Repayment Behaviour decreased (11,466 → 10,666)	Borrowers in Group B are repaying slightly less principal on average, which may indicate weaker repayment discipline or shorter repayment progress.
Default Risk decreased (recoveries 138.9 → 119.2)	Lower recoveries suggest fewer defaults or lower recovery activity, implying the credit risk profile may have improved slightly.

Key Insight

Group B shows:

- Higher loan adoption
- Higher borrowing amounts
- Slightly weaker repayment levels
- Lower default recoveries

Net effect: loan growth improved without a visible increase in default risk.

Step 8 — Business Decision

Based on the observed metrics Strategic Recommendation:

We should adopt the 2013 strategy cautiously.

Reasoning:

- Loan acceptance increased → market demand improved
- Average loan size increased → higher revenue potential
- Default risk decreased → no immediate deterioration in credit quality
- **Repayment behaviour slightly declined** → requires monitoring

Recommended Business Action

1. Adopt the lending strategy used in 2013.
2. Implement monitoring controls:
 - Track repayment trends for the next lending cycle.
 - Segment borrowers by risk grade or credit score.
3. Risk-adjusted expansion strategy:
 - Offer larger loans or favourable rates primarily to low-risk borrowers.
 - Apply stricter evaluation for high-risk segments.

Executive Summary

The 2013 lending environment appears to increase loan adoption and loan size while keeping default risk stable or slightly improved.

However, the drop in repayment behaviour signals a potential early warning, so scaling the strategy should be accompanied by risk monitoring and borrower segmentation.

DBMS & Data Architecture

ACID Properties:

1. A user initiates a fund transfer from Account A to Account B, but the app crashes midway. How would you ensure that no partial transaction corrupts the account balances?

Answer:

- **Atomicity:** In KreditBee, a loan disbursement must be treated as a single, indivisible transaction. Either both actions happen (debiting KreditBee's escrow account and crediting the user's bank account), or neither happens. If the app crashes midway, the database automatically rolls back the entire transaction, ensuring no funds are lost in transit.
- **Durability:** Once the disbursement is successful, the transaction is committed to the database permanently. Even if the database server crashes one second later, the record of that successful transfer remains intact upon reboot.

2. Two traders place orders for the same stock at nearly the same time. How would you ensure that each transaction executes reliably without interference?

Answer:

- **Isolation:** The database must use transaction isolation levels (like Read Committed or Serializable). If a user accidentally taps the "Pay EMI" button twice rapidly, the database processes the first

transaction independently and locks the row. The second transaction will either wait or fail, preventing the user from being double-charged.

- **Consistency:** The system enforces database rules so that the user's outstanding loan balance never drops below zero incorrectly, maintaining the integrity of the ledger.

Indexing:

1. You need to generate real-time dashboards showing account balances and transactions for high-net-worth clients. How can indexing help here?

Answer:

Without indexing, the database would have to perform a full table scan across millions of KreditBee users to find one person's data. By creating a **B-Tree index** on frequently queried columns like **userID** and **accountID**, the database can locate the user's transaction history in milliseconds, allowing the real-time dashboard to load instantly.

2. Transaction data is frequently updated and queried by transaction ID and date. What indexing strategies would you use to optimize performance without increasing update costs too much?

Answer:

Indexing helps by creating a B-Tree index on frequently queried columns like `userID` and `accountID`. Instead of forcing the database to perform a slow, full table scan across millions of KreditBee users to find one person's data, the index allows the database to locate the user's transaction history in milliseconds, ensuring the real-time dashboard loads instantly.

Partitioning

1. Your transaction table contains billions of records across multiple years. How would you design it to allow efficient querying and archiving?

Answer:

We would implement **Time-Based (Range) Partitioning**. The `Repayment_Schedule` and `Transactions` tables would be partitioned by month or year (e.g., `transactions_2024_Q1`, `transactions_2024_Q2`). Active queries usually fetch recent data, so the database only scans the most recent partition. Older partitions (like 2021 data) can be compressed and archived to cheaper storage without affecting current application performance.

2. You need to comply with data residency laws that require customer data to stay within a particular region. How can partitioning help?

Answer:

We would use **Geo-Partitioning (List Partitioning)** based on location. To comply with Reserve Bank of India (RBI) data localization mandates, all transactions originating from Indian users are routed to and stored strictly within the "India" partition on local servers (e.g., AWS ap-south-1).

Sharding

1. During peak hours, your database hits performance limits while processing millions of microtransactions. How would you use sharding to scale horizontally?

Answer:

During high-traffic events (like festive season sales driving up "Flexi Loan" requests), a single database server becomes a bottleneck. We would implement **Horizontal Sharding** using **userID** as the Shard Key.

- User IDs 1 to 5,000,000 are stored on Server A.
- User IDs 5,000,001 to 10,000,000 are stored on Server B. This distributes the read/write load evenly across multiple machines, scaling horizontally to handle unlimited microtransactions.

2. You've implemented sharding by customer ID, but now need to perform analytics on global transaction patterns. What problems might you face?

Answer:

Sharding optimizes localized operations but breaks global analytics. To calculate the "Total Daily Default Rate" across all users, the system would have to perform **Cross-Shard Joins**, querying Server A, B, and C simultaneously, which is incredibly slow and resource-heavy. *Solution:* Analytical queries should not be run on the sharded operational database (OLTP). Instead, the data from all shards should be streamed into a centralized Data Warehouse for analytics.

Normalization and Denormalization

1. You're designing a schema to store customer profiles, accounts, and KYC documents. How would normalization help here?

Answer: Normalization reduces data redundancy and prevents update anomalies. By separating the data into distinct tables (**User_Profile**, **KYC_Details**, **Loan_Account**) linked by Foreign Keys, we ensure that if a user updates their phone number, it only needs to be updated in one single place (**User_Profile**). It ensures strict data integrity for compliance purposes.

2. To speed up monthly financial reporting, would you prefer a normalized or denormalized schema? Explain why.

Answer: A **Denormalized Schema** is preferred for reporting. Generating monthly financial reports on a highly normalized database requires complex, slow **JOIN** operations across a dozen tables. A denormalized schema flattens this data into wide summary tables, accepting some data redundancy in exchange for massive improvements in read-heavy analytical performance.

Data Lakes vs. Data Warehouses

1. You receive unstructured logs, real-time clickstream data, and structured account data. Where would you store each, and why?

Answer:

- **Data Lake (e.g., AWS S3):** Used for the unstructured server logs and real-time app clickstream data. Lakes handle raw, unformatted data cheaply and efficiently at massive scale.
- **Data Warehouse (e.g., Snowflake, BigQuery):** Used for the structured account and loan data. Warehouses require a predefined schema and are optimized for fast, complex SQL querying by data analysts.

2. How would a data warehouse support fraud detection and risk analysis in a FinTech company?

Answer: A Data Warehouse excels at **OLAP (Online Analytical Processing)**. It can instantly aggregate years of historical loan data. Risk analysts can write queries to detect complex fraud patterns—such as identifying if multiple defaulted loans across different states share the same underlying device ID or IP address over a 24-month period.

Data Ingestion and Pipelines

1. Describe how you would ingest real-time payments data from mobile apps into your system for fraud monitoring.

Answer: We would use a **Streaming Data Pipeline** powered by Apache Kafka or AWS Kinesis. As soon as a user makes a payment via KreditBee UPI, the event is pushed into a Kafka topic in real-time. A fraud-detection machine learning model consumes this stream instantly, analyzing the transaction context before it even hits the main database.

2. You need to import credit score data from third-party providers once a day. What ingestion approach would you use?

Answer: We would use a **Batch Data Ingestion** strategy orchestrated by a scheduling tool like Apache Airflow. At 2:00 AM during low-traffic hours, an automated ETL job triggers, calls the TransUnion/CIBIL APIs, fetches the updated credit scores in bulk and updates the **User_Profile** table.

ETL (Extract, Transform, Load)

1. How would you design an ETL pipeline to consolidate customer investment portfolios from multiple banking systems into a central analytics platform?

Answer:

- **Extract:** Pull raw data incrementally from various partner bank APIs.
- **Transform:** Cleanse the data. This involves standardizing currency formats, resolving mismatched dates (e.g., MM-DD-YYYY vs DD-MM-YYYY), and deduplicating records.
- **Load:** Insert the transformed, clean data into a centralized Data Warehouse schema, ready for business intelligence dashboards.

2. An ETL job failed after partial data load. How would you maintain consistency and resume safely?

Answer: The ETL pipeline must be built with **Idempotency** and **Checkpointing**. If a job fails midway, it should not blindly restart from the beginning (which would create duplicate loan entries). Instead, it reads the last successful checkpoint. Furthermore, "Upsert" logic (**INSERT ON CONFLICT UPDATE**) ensures that if a row was partially loaded, running the job again simply overwrites it safely rather than duplicating it.

Additional DBMS Topics

1. Regulators request a complete audit trail of customer loan applications. How would you ensure your DBMS can support this?

Answer: We would implement **Audit Logging** with an append-only architecture. Instead of just updating a user's loan status from "Pending" to "Approved," we write a new timestamped row to an **Application_History** ledger table. This provides an immutable, version-controlled timeline of exactly when decisions were made, satisfying regulatory compliance.

2. Your machine learning model needs daily updated training data on loan defaults. How would you manage the data pipeline to feed the model efficiently?

Answer: We would use **Change Data Capture (CDC)** combined with Incremental ETL. Instead of reloading the entire 10-year loan history every day, CDC listens to the database's transaction logs and extracts *only* the rows that changed status to "Defaulted" in the last 24 hours. This snapshot is efficiently fed into the ML feature store.

3. How would you detect and resolve duplicate records in credit card applications ingested from multiple sources?

Answer: During the "Transform" phase of the ETL pipeline, we apply **Fuzzy Matching algorithms** and deterministic data quality checks. We would compare composite keys (e.g., **PAN Card Number + Date of Birth**). If exact matches fail due to typos, algorithms like Levenshtein distance can flag similar names (e.g., "Rahul Sharma" vs "Rahul Sarma") for manual resolution.

4. You are tasked with maintaining real-time analytics dashboards for customer transactions, credit scores, and investments. What DBMS architecture would you propose?

Answer: We would deploy a **Hybrid Transactional/ Analytical Processing (HTAP)** architecture.

- The core application runs on an OLTP database (like PostgreSQL).
- We create **Materialized Views** to pre-calculate heavy aggregations (like "Total Disbursals Today") so the dashboard doesn't re-calculate the math on every page load.
- A caching layer (like **Redis**) sits in front to serve these dashboard metrics instantly, taking the heavy read load off the primary database.

Week II — SQL ETL, Excel EDA & Advanced Analytics

Cover Page

- **Project Title:** Financial Risk and Performance Analysis in Consumer Lending
- **Dataset Name:** loan_final313_.csv

Section 1: Project Overview & Objectives

Project Overview:

This project focuses on the end-to-end data analytics process in the consumer lending landscape. Using the loan_final313_.csv dataset, we performed an ETL (Extract, Transform, Load) process using Colab to clean and structure the data. The transformed dataset was then analyzed in Excel to uncover patterns in borrower demographics, loan characteristics, and repayment histories.

Objectives:

- Apply advanced data analysis techniques to deeply understand the loan portfolio.
- Assess loan risk profiles and evaluate financial performance metrics.
- Develop an interactive Excel dashboard for real-time strategic decision-making.
- Provide strategic recommendations for optimizing loan portfolio management and mitigating default risks.

Section 2: ETL Phase (SQL)

(Team Note: This section covers all of Task 3 from the instructions.)

SQL Environment Setup & Queries Link:

- [Code](#)

Summary of ETL Execution:

- **Extract:** Created the loans table with appropriate data types and successfully loaded the raw .csv file. Verified record counts and unique categorical values.
- **Transform:** Replaced missing annual_inc values with the median. Standardized home_ownership to uppercase. Engineered new columns: profitability (Total Payment - Loan Amount) and risk_flag (1 for Bad Loan, 0 for Good Loan).
- **Load:** Created a final loans_cleaned table. Added default_rate_indicator and income_to_loan_ratio, extracted numeric loan terms, and exported the clean dataset to a new CSV for Excel analysis.

Section 3: Excel Analysis

(Team Note: This section covers all 20 questions from Task 4 and the Dashboard creation.)

Excel File Link:

- [Cleaned csv](#)

Docs File Link:

- [Task - 4 \(EDA\)](#)


Summary of Exploratory Data Analysis (EDA):

- **Distribution & Trends:** Analyzed the distribution of loan amounts, interest rate trends over the years, and the correlation between income categories and loan amounts.
- **Risk & Performance:** Calculated year-over-year default rates, recovery rates across loan grades, and grade-wise profitability.
- **Advanced Computation:** Utilized SUMPRODUCT and SUMIF to calculate weighted average interest rates, and combined IF, MODE, and MATCH to determine the most common loan purpose for defaulted loans by region.
- **Visual Formatting:** Applied conditional formatting to highlight high-risk loans (high interest rates + low grades).

Interactive Dashboard:

(Final Excel Dashboard here)

- Dashboard



Section 4: Insights & Observations

Top 5 Findings:

1. **Income vs. Risk:** Our predictive probability model revealed that applicants with lower annual incomes (<\$80,000) requesting higher loan amounts (>\$10,000) carry a specific default likelihood of 12.67%. Furthermore, volume-adjusted metrics showed Medium-Income borrowers actually yield the highest weighted interest (14.28%), indicating they take on larger high-rate loans than the Low-Income bracket.
2. **Employment Stability:** The analysis confirmed that shorter employment histories (under 5 years), when combined with high loan requests and lower income, create a compounded high-risk profile for default.
3. **Loan Purpose Trends:** Advanced array modeling definitively proved that "debt_consolidation" is the absolute primary driver of bad loans. Borrowers taking loans to pay off existing debt represent the highest default frequency by a massive margin.
4. **Interest Rate Impact:** High-interest-rate loans naturally correlated with higher default probabilities. The data shows that the safest borrowers (High Income: 100k-150k) secure the lowest weighted average interest rates (13.93%), while riskier brackets are forced to absorb higher rates.
5. **Regional Risk:** Regional analysis (spanning Cannught, Munster, Ulster, Northern-Irl, and Leinster) revealed a uniform risk factor: geography matters far less than loan purpose. Debt consolidation remains the leading cause of default across *every single territory*, indicating systemic borrower behavior risk rather than localized economic issues.

Recommended Business Actions:

- **Stricter Screening:** Implement rigid income-to-loan-amount and employment-length thresholds before approving high-value loans, specifically for the Medium and Low-income brackets.
- **Optimize Pricing:** Adjust interest rates based on the borrower's Debt-to-Income (DTI) ratio rather than relying solely on high interest rates to offset bad risk.
- **Monitor High-Risk Purposes:** Cap total loan amounts for high-risk categories like Debt Consolidation unless the borrower has a flawless historical repayment record.

Section 5: Reflection

Challenges Faced:

- **SQL Transformation:** Managing null values and ensuring categorical text standardizations (like converting to uppercase and standardizing region names) did not break subsequent query logic during the initial extraction phase.
- **Excel Dashboard Connectivity:** Discovered that separate pivot tables create individual pivot caches, causing interactive slicers to fail. The dashboard architecture had to be optimized to run from a unified data model.
- **Advanced Excel Formulas:** Finding the most frequent loan purpose for defaulted loans required nesting INDEX, MATCH, IF, and MODE functions inside an ARRAYFORMULA. This was highly complex because the standard MODE function does not natively support text strings, requiring a numeric translation array.

Learning Outcomes:

- **Mastered the end-to-end data pipeline:** Successfully executed a full analytics workflow from raw SQL extraction and data engineering to front-end visualization and dynamic dashboarding.
- **Financial Acumen:** Gained a deep, practical understanding of financial risk metrics, including volume-adjusted weighted interest rates, conditional probability modeling, and default rate analysis.
- **Advanced Computational Logic:** Improved advanced Excel skills, specifically mastering array-based statistical computations, conditional aggregation (AVERAGEIFS, SUMPRODUCT), and dynamic data extraction.

Conclusion


This project significantly enhanced our technical, analytical, and problem-solving abilities. It strengthened our confidence in managing a complete analytics workflow — from SQL-based data preparation to interactive dashboard creation and strategic recommendation development.

Overcoming practical challenges, particularly in dashboard structuring and slicer integration, was a valuable learning experience that deepened our understanding of real-world data analytics implementation.




Week III — PowerBI, Automation & GenAI

Phase 1: Power BI Dashboarding

This phase is about visual storytelling. We will use cleaned `loans_cleaned` dataset and build an interactive dashboard that answers the business questions.





Lending Trends & Growth




Risk Analysis & Loan Status


grade: region: year: annual_income: emp_length:



Risk Matrix: Default Rate by Grade & Tenure

grade	0.50	1.00	2.00	3.00	4.00	5.00	6.00	6.05	7.00	8.00	9.00	10.00	Total
A	0.13	0.13	0.13	0.12	0.13	0.13	0.13	0.13	0.13	0.13	0.13	0.13	0.13
B	0.13	0.13	0.13	0.13	0.13	0.14	0.14	0.13	0.13	0.13	0.13	0.13	0.13
C	0.13	0.13	0.13	0.13	0.13	0.13	0.13	0.12	0.13	0.13	0.12	0.12	0.13
D	0.13	0.13	0.13	0.13	0.13	0.13	0.13	0.12	0.13	0.13	0.13	0.12	0.13
E	0.13	0.13	0.13	0.13	0.13	0.13	0.13	0.11	0.13	0.12	0.12	0.12	0.13
F	0.13	0.14	0.13	0.13	0.13	0.14	0.14	0.12	0.13	0.13	0.13	0.13	0.13
G	0.14	0.15	0.12	0.13	0.13	0.15	0.14	0.14	0.14	0.13	0.13	0.13	0.13
Total	0.13	0.13	0.13	0.13	0.13	0.13	0.13	0.12	0.13	0.13	0.13	0.13	0.13






PowerBI Folder Github Link:

[Lending-Loan-Analysis-PowerBI](#)

Github Link:

[Loan Risk Automation Gen AI Insights](#)



Key Insights

Unified Risk Visibility: The integration of Power BI and n8n transforms raw financial data into a continuous intelligence loop. Power BI provides macro-level portfolio monitoring, while n8n handles micro-level, event-driven automated reporting.

Reduction in Analytical Latency: By replacing manual data extraction with an automated n8n pipeline, the time from data ingestion to executive reporting (via GenAI and Gmail APIs) is reduced from hours to seconds.

Enhanced Borrower Transparency: Utilizing LLaMA 3.3 70B via webhook triggers allows for complex financial logic (such as Debt-to-Income impact and Grade classifications) to be translated into empathetic, easily understandable language for customers.

Scalable Architecture: The use of REST APIs, webhook triggers, and automated payload formatting ensures this solution can easily integrate with enterprise CRMs, core banking systems, and real-time dashboards.

Risk Insights & Business Actions

1. Insight: High Correlation Between DTI, Low Income, and Default Rates

Risk Factor: The Power BI dashboard and AI analysis consistently flag borrowers with Debt-to-Income (DTI) ratios above 40% and lower income brackets as the primary drivers of "Bad Loans."

Business Action: Tighten automated underwriting thresholds for Grade E and F loans. Implement dynamic interest rate pricing for applicants crossing the 35% DTI mark to offset potential default risks.

2. Insight: Executive Reporting Bottlenecks

Risk Factor: Relying on manual end-of-month reporting leaves leadership blind to rapid shifts in portfolio health.

Business Action: Deploy Workflow 1 to act as an automated "AI Chief Risk Officer." Schedule daily or weekly GenAI summaries pushed directly to executive inboxes, highlighting immediate risk drivers without requiring manual dashboard analysis.

3. Insight: Borrower Confusion Regarding Rejections

Risk Factor: Lack of transparency in loan status classifications leads to poor customer experience and increased support ticket volume.

Business Action: Implement Workflow 2 into the customer portal. Trigger an on-demand webhook that instantly generates and emails a personalized, jargon-free explanation to the borrower, comparing their profile to a benchmarked "Good Loan."

Conclusion

This end-to-end data product successfully bridges the gap between static analytics and active business operations. While the Power BI dashboard serves as the central hub for discovering deep quantitative risk trends, the n8n + GenAI architecture acts as the operational engine—automating both executive-level reporting and customer-facing communication. By treating data not just as numbers to be visualized, but as triggers for automated AI logic, this solution significantly reduces manual overhead, accelerates strategic decision-making and establishes a modern, transparent lending pipeline.